

Detection Engineering for Data Loss Prevention: Designing Splunk-Based Alerts and Response Plans for Data Exfiltration Incidents

Pedram Shahbazi, Robert (Ash) Drake,

With collaboration of Samsung Austin Semiconductor (Engineer Eric Phillpott)

Project (CS-6399-D10)

Dr. Erdoğan Doğdu, Dr. Roya Choupani

Master of Computer Science

Department of Computer Science

Angelo State University

Spring 2026

Project Description

Data exfiltration—the unauthorized transfer of sensitive information outside an organization’s-controlled environment—is one of the most critical threats facing modern enterprises. It frequently occurs through seemingly legitimate channels such as web uploads, cloud storage services, email, and DNS, and is often carried out by compromised accounts or malicious insiders attempting to evade traditional security controls. Conventional Data Loss Prevention (DLP) solutions, while effective at policy-based blocking and content inspection, often operate as isolated components that generate high volumes of uncorrelated alerts, leading to alert fatigue and missed incidents.

Security Information and Event Management (SIEM) platforms, such as Splunk, aggregate telemetry from diverse sources including network devices, endpoints, SaaS platforms, and DLP tools. By working on aggregated data within such a platform, it becomes possible to engineer higher-fidelity detections that capture multi-step exfiltration behaviors and contextual signals that single-point tools cannot provide. This project aims to systematically apply detection-engineering principles to DLP-relevant scenarios using Splunk, focusing on the design of correlation-based alerts and associated incident response plans for data exfiltration events.

This project will design, implement, and evaluate a set of Splunk-based detection use cases covering multiple exfiltration vectors (e.g., large outbound transfers, exfiltration to unsanctioned cloud services, email-based exfiltration, and DNS-mediated exfiltration). For each use case, the project will define the required data sources, implement detection logic using Splunk Processing Language (SPL), and specify structured response playbooks. Evaluation will be carried out using synthetic or lab-generated datasets that emulate both benign activity and exfiltration scenarios, with emphasis on detection coverage, false-positive behavior, and operational usefulness for a Security Operations Center (SOC).

Related Work

A significant body of research has addressed data exfiltration from both theoretical and practical perspectives. Academic studies have proposed anomaly-detection techniques over network flows, host activity, and application logs, often relying on statistical models or machine learning to identify deviations from normal behavior. These works demonstrate that volume-based and behavior-based features, such as sudden spikes in outbound bandwidth, unusual communication patterns, or abnormal DNS usage, can provide useful indicators of exfiltration activity. However, many of these approaches are evaluated in controlled settings and are not directly integrated into operational SIEM workflows.

In parallel, commercial DLP products and CASB solutions provide policy-based inspection and blocking capabilities across endpoints, email, web, and cloud applications. They excel at enforcing predefined rules (e.g., preventing specific document types from leaving the environment) and at identifying content that matches sensitive data classifications.

Nonetheless, these tools are constrained by their local perspective: they typically see only the channel they control and can generate numerous stand-alone alerts that are difficult to triage without broader context. Moreover, sophisticated attackers can adapt their behavior to remain below thresholds or to use channels that are less strictly monitored.

Modern SIEM and SOAR (Security Orchestration, Automation, and Response) platforms attempt to bridge this gap by aggregating telemetry and providing detection content and playbooks for common threats, including data exfiltration. Vendor-supplied “out-of-the-box” rules and analytics often target patterns such as anomalous outbound traffic volumes, suspicious cloud service usage, or DNS tunneling signatures. While useful, these generic detections must be tuned to local environments and may not reflect specific organizational risk scenarios or data flows. Additionally, many SOC playbooks for exfiltration focus primarily on response orchestration and ticketing, leaving the underlying detection logic implicit.

This project distinguishes itself from prior work and existing products in several ways. First, it treats detection engineering for DLP scenarios as a structured, end-to-end process: from threat modeling and use-case definition, through data modeling and feature selection, to SPL implementation and iterative tuning. Second, it explicitly focuses on leveraging aggregated multi-source data within Splunk, aiming to correlate network, DNS, SaaS, and DLP events to increase signal fidelity and reduce false positives. Third, it couples each detection with a concrete incident response plan, thereby closing the gap between “alert fired” and “incident handled.” In doing so, the project contributes a practical framework and set of artefacts that can be adapted by organizations seeking to operationalize data-exfiltration detection using Splunk.

Requirements

Functional Requirements

The proposed system shall:

1. Ingest and normalize security-relevant telemetry into a Splunk environment, including:
 - Network and/or proxy logs containing source, destination, and byte-count information.
 - DNS logs capturing query names, types, and client information.

- SaaS or office-suite audit logs (e.g., email events, file sharing activities, mailbox exports).
 - DLP or security-tool logs indicating policy violations or relevant alerts, where available.
2. Define and implement a set of at least four data-exfiltration detection use cases, such as:
 - High-volume outbound data transfers from an internal host or user.
 - Data transfers to unsanctioned or rare external destinations, including cloud storage services.
 - Email-based exfiltration, including mass forwarding, bulk attachment sending, or mailbox/PST exports.
 - DNS-based exfiltration patterns, characterized by abnormal query volume, length, or entropy.
 3. For each use case, implement SPL-based detection logic that:
 - Aggregates and correlates relevant events over configurable time windows.
 - Computes pertinent features (e.g., total bytes, number of events, rarity scores, deviation from baseline).
 - Generates alerts or notable events with sufficient contextual enrichment to support investigation.
 4. Design and document an incident response playbook for each detection use case, articulating:
 - Initial triage steps and verification checks.
 - Enrichment queries and context gathering (e.g., user/asset history, previous alerts).
 - Containment, eradication, and recovery actions, as well as communication and escalation pathways.
 5. Provide an evaluation of the detections and playbooks, including:
 - Results of applying the detections to synthetic or lab-generated exfiltration scenarios.
 - Observations on false positives when applied to benign data.
 - A qualitative assessment of the operational suitability of the alerts and response procedures.

Non-Functional Requirements

The system shall satisfy the following non-functional properties:

- **Performance:** Detection searches must complete within acceptable execution times for typical query horizons (e.g., daily or hourly windows) on the target data volume.
- **Scalability:** Detection logic should be designed using efficient aggregations, avoiding unnecessary complexity that would prevent scaling to larger datasets.
- **Explainability:** Each detection shall be accompanied by clear documentation of its rationale, thresholds, assumptions, and potential blind spots.
- **Maintainability:** SPL rules and playbooks shall be modular and parameterized, enabling tuning and extension without extensive rework.
- **Security and Privacy:** Use of any real data must comply with organizational and legal privacy requirements; otherwise, anonymized or synthetic data will be employed.

Methods/Tools

The project will follow an iterative detection-engineering methodology comprising the following steps: threat modeling, data exploration, feature design, detection implementation, tuning, and response-playbook definition. For each exfiltration scenario, the process will begin by articulating an attacker hypothesis (e.g., “an insider compresses sensitive files and uploads them to an unsanctioned cloud storage site”), mapping it to observable signals in available logs, and enumerating the features that can distinguish this behavior from benign activity.

Splunk will serve as the central analytics platform. Its search and correlation capabilities (via SPL) will be used to implement feature extraction and detection logic over aggregated data. Saved searches or scheduled searches will implement the detections, while dashboards or simple views will assist in exploration and evaluation. Where appropriate, baseline and anomaly-detection approaches will be implemented using SPL constructs (e.g., rolling averages, standard deviations, or percentile-based thresholds). If time permits, basic machine-learning models may be integrated using Splunk’s machine-learning enhancements or external scripts, though the core emphasis will remain on interpretable rule-based detections.

For response modeling, the project will define playbooks in a platform-agnostic manner, specifying sequences of actions and decision points that align with SOC practices. If a suitable orchestration platform is available, selected playbooks may be implemented as automated or semi-automated workflows; otherwise, they will be represented through flow diagrams and procedural documentation. Auxiliary tools, such as scripting environments

for data generation and version-control systems for managing SPL and configuration artefacts, will support repeatability and disciplined development.

Project Timeline

The work will be organized into phased milestones over the project period:

- **Weeks 1–2:** Finalize problem definition, scope, and detailed use cases; perform focused literature and industry review; prepare initial project environment.
- **Weeks 3–4:** Onboard and normalize initial data sources into Splunk; define the conceptual data model and key fields; design feature sets for the selected use cases.
- **Weeks 5–7:** Implement detection logic for core use cases (high-volume outbound transfers, unsanctioned destinations, email-based exfiltration); perform initial validation and adjustments.
- **Weeks 8–9:** Implement DNS-based exfiltration detection and any additional high-priority use cases; draft corresponding incident response playbooks.
- **Weeks 10–11:** Design and run test scenarios for benign and malicious behaviors; tune detections; refine playbooks based on findings.
- **Weeks 12–14:** Consolidate results; complete the final report and documentation; prepare and rehearse the project presentation.

Dependencies between tasks will be managed to ensure that later phases (e.g., evaluation) can be carried out even if some integrations are only partially available, by relying on synthetic or constrained datasets where necessary.

References

- Splunk Threat Research Team. “Data Exfiltration Detections: Threat Research Release, June 2021.” 2021.
- Splunk. “Detecting Data Exfiltration Activities.” Splunk Lantern.
- Splunk. “Machine Learning in Security: Detect DNS Data Exfiltration Using Deep Learning.” 2022.
- Splunk. “Machine Learning in Splunk Enterprise Security: Unleashing Hidden Threats.” 2025.
- SearchInform. “SIEM Workflow Automation: Streamlining Incident Response.” 2025.
- Radiant Security. “SOC Playbook Examples for Real-World Cyber Threats.” 2026.
- Hamdan, M. “Data Exfiltration Explained: How to Use Splunk and Wireshark to Detect It.” 2025.
- Splunk Lantern. “Data Loss Prevention Data.” 2025.

- PuppyGraph. “How to Build a SIEM System: Architecture & Tools.” 2026.